Rationale

I built this because so many outlets publish the same news. Forty headlines can be forty policies or one policy getting approved, and a cancellation a few days later reads almost like the original. The goal is to count policies, put them out for public review so anyone can see which states are more proactive on AI, and later follow how much of each policy gets implemented. Every policy keeps the news reports behind it. Every dropped item keeps its reason.

Sources. I used Google News RSS as the main source because it is free, needs no key and can be searched by state and by month. For each state and month I ran four searches: the government itself, policy instruments, public services, and infrastructure and rules. That came to 3,068 searches, and every response is saved so the dataset can be rebuilt offline. GDELT and Bing added a little, and GDELT was rate limited on most calls. I wrote code for three paid news APIs but never ran it with a key, so I do not count them.

States. I ran one common search across all 28 states and picked the working set from how much came back and how the states are spread across the country. Three were added to compare against a separate list built from official documents. The thirteen are Andhra Pradesh, Assam, Bihar, Gujarat, Karnataka, Kerala, Madhya Pradesh, Maharashtra, Odisha, Rajasthan, Tamil Nadu, Telangana and Uttar Pradesh. I searched English news only, so a low count can simply mean a state's news is in another language.

Collecting and cleaning. The state is read from the headline, never from the search. The list of names covers spellings like UP and Uttar-Pradesh, common misspellings, and short forms like MP that count only next to a government word. Dates are converted to India time, and an item with no date is dropped, never filled in.

The main algorithm. Removing duplicates is the core of the work. Headlines from the same state, at most five days apart, get a score from the rare words they share, and pairs above the bar are merged. A cancellation never merges with the original. The final check found two such wrong merges, and fixing them took the total from 714 to 716 policies, with 1,306 news reports behind them. Inside each group, the headline with the highest average score represents the policy, and every policy carries a confidence level and the number of outlets that reported it.

Presentation. I kept the dashboard simple: big numbers and a growth chart first, then filters and the list, with every policy linked to its news reports. The charts show what was reported. They cannot tell me whether a policy was delivered.

What is weak. An AI assistant screened the headlines and labelled the pairs that set the bar. I have not read every article, and the model check inside the pipeline is written but not yet run. The next thing I would do is link each announcement to an order, a tender or a budget line. That is where tracking implementation starts.